

Machine Learning

Introduction

Welcome

Apple - iPhoto - New full-Sc X

www.apple.com/ilife/iphoto/

Store Mac iPod iPhone iPad iTunes Support

iLife '11

iPhoto iMovie GarageBand Video Showcase Resources Upgrade Now

 iPhoto '11

From your Facebook Wall to your coffee table to your best friend's inbox (or mailbox). Do more with your photos than you ever thought possible. And do it all in one place. iPhoto.

 Watch the iPhoto video ▶



What's New in iPhoto What is iPhoto?





Machine Learning

- Grew out of work in AI
- New capability for computers

Examples:

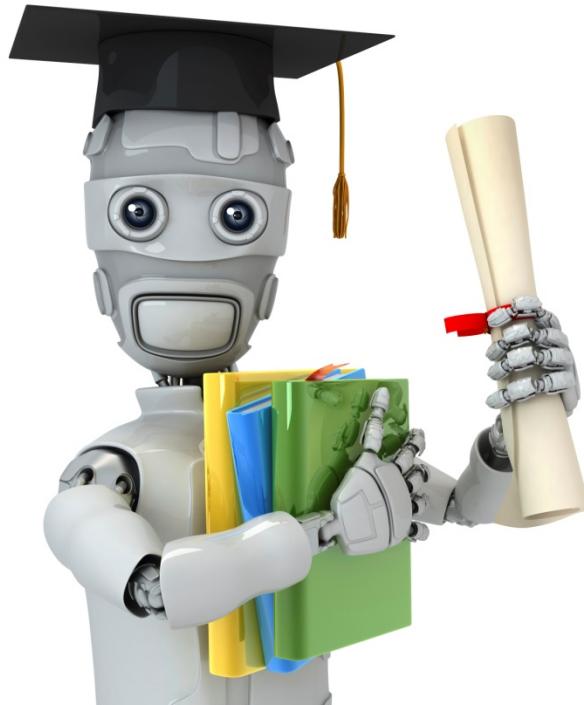
- Database mining
 - Large datasets from growth of automation/web.
E.g., Web click data, medical records, biology, engineering
- Applications can't program by hand.
 - E.g., Autonomous helicopter, handwriting recognition, most of Natural Language Processing (NLP), Computer Vision.

Machine Learning

- Grew out of work in AI
- New capability for computers

Examples:

- Database mining
 - Large datasets from growth of automation/web.
E.g., Web click data, medical records, biology, engineering
- Applications can't program by hand.
 - E.g., Autonomous helicopter, handwriting recognition, most of Natural Language Processing (NLP), Computer Vision.
- Self-customizing programs
 - E.g., Amazon, Netflix product recommendations
- Understanding human learning (brain, real AI).



Machine Learning

Introduction

What is machine learning

Machine Learning definition

- Arthur Samuel (1959). Machine Learning: Field of study that gives computers the ability to learn without being explicitly programmed.
- Tom Mitchell (1998) Well-posed Learning Problem: A computer program is said to *learn* from experience E with respect to some task T and some performance measure P, if its performance on T, as measured by P, improves with experience E.

“A computer program is said to *learn from experience E* with respect to some task T and some performance measure P, if its performance on T, as measured by P, improves with experience E.”

Suppose your email program watches which emails you do or do not mark as spam, and based on that learns how to better filter spam. What is the task T in this setting?

- Classifying emails as spam or not spam. T ←
- Watching you label emails as spam or not spam. E ←
- The number (or fraction) of emails correctly classified as spam/not spam. P ←
- None of the above—this is not a machine learning problem.

Machine learning algorithms:

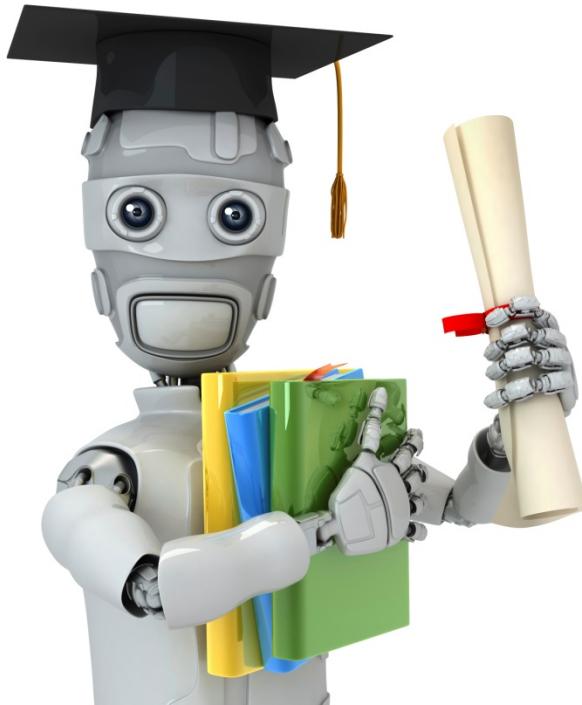
- Supervised learning
- Unsupervised learning



Others: Reinforcement learning, recommender systems.

Also talk about: Practical advice for applying learning algorithms.



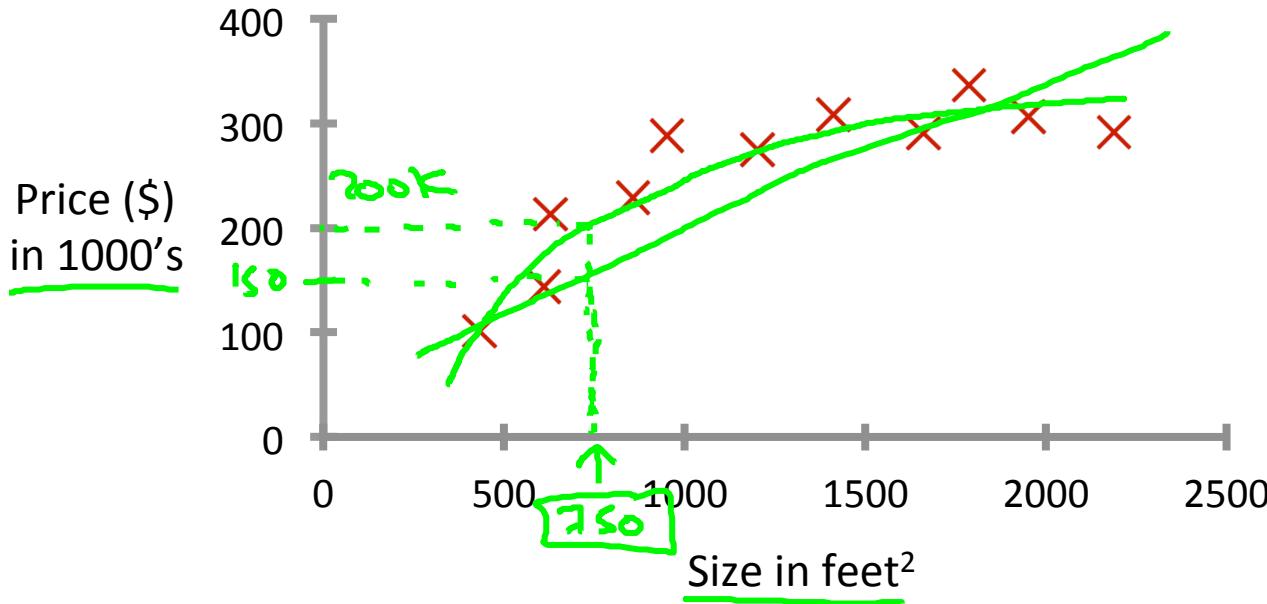


Machine Learning

Introduction

Supervised Learning

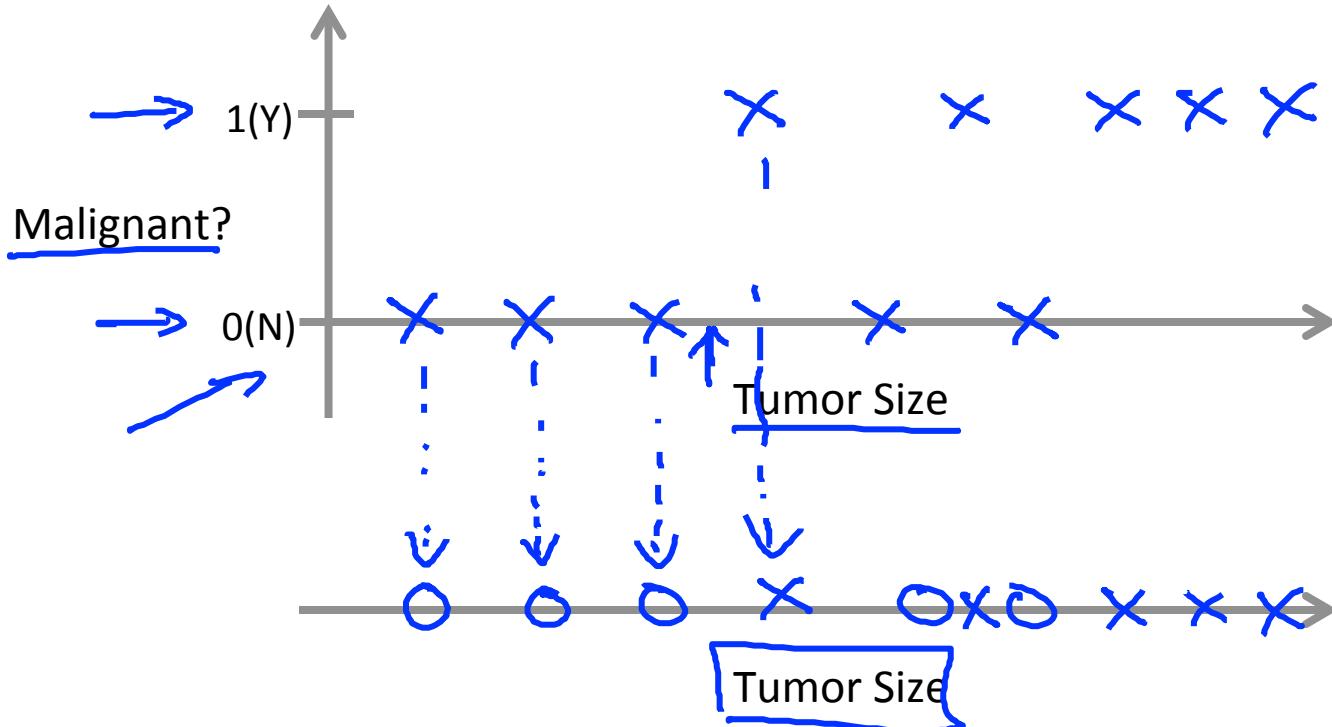
Housing price prediction.



Supervised Learning
'right answers' given

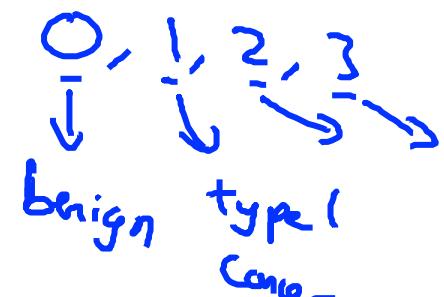
Regression: Predict continuous valued output (price)

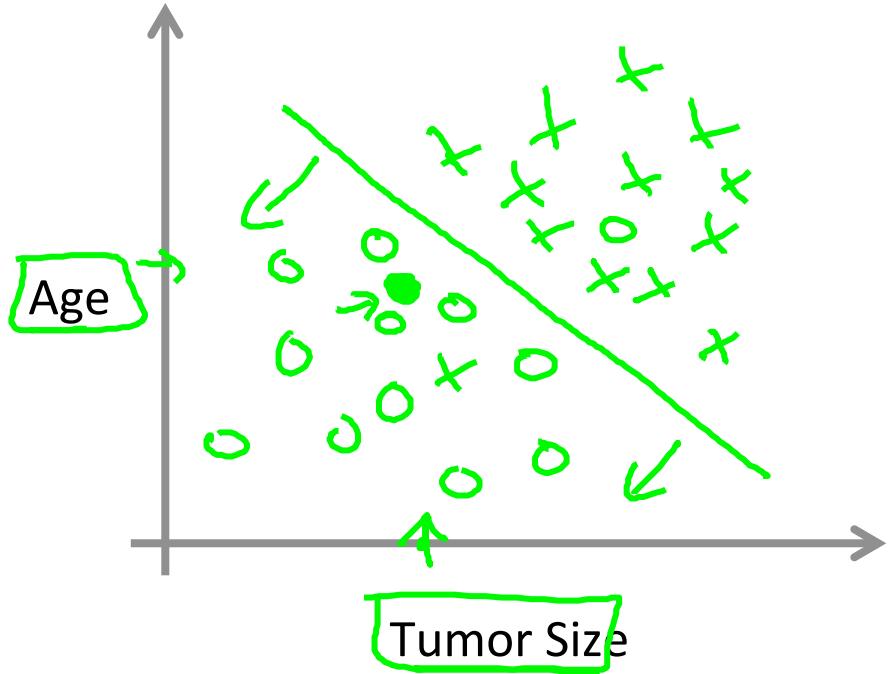
Breast cancer (malignant, benign)



Classification

Discrete valued output (0 or 1)





You're running a company, and you want to develop learning algorithms to address each of two problems.

1000's

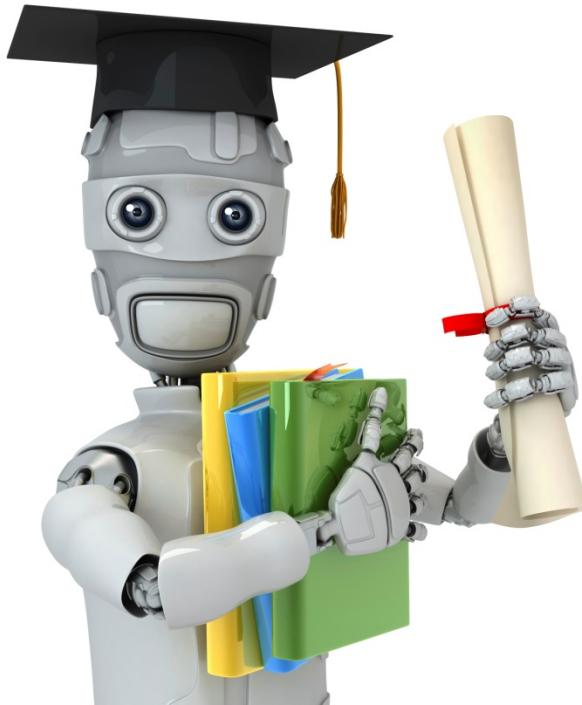
→ Problem 1: You have a large inventory of identical items. You want to predict how many of these items will sell over the next 3 months.

→ Problem 2: You'd like software to examine individual customer accounts, and for each account decide if it has been hacked/compromised.

→ 0 - not hacked
→ 1 - hacked

Should you treat these as classification or as regression problems?

- Treat both as classification problems.
- Treat problem 1 as a classification problem, problem 2 as a regression problem.
- Treat problem 1 as a regression problem, problem 2 as a classification problem.
- Treat both as regression problems.

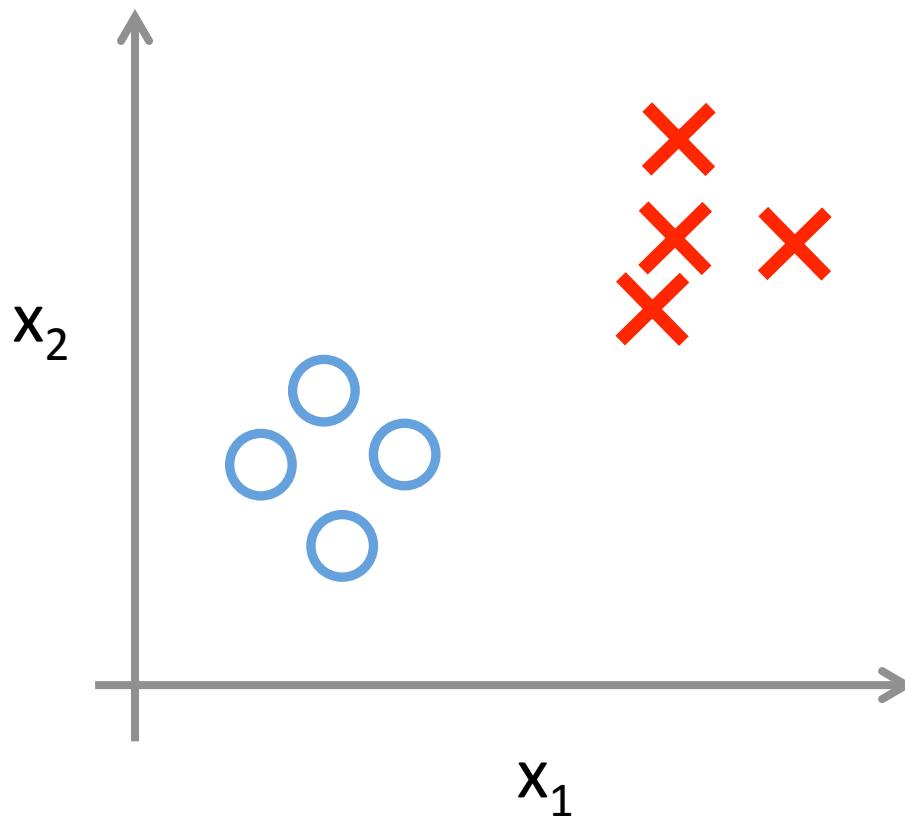


Machine Learning

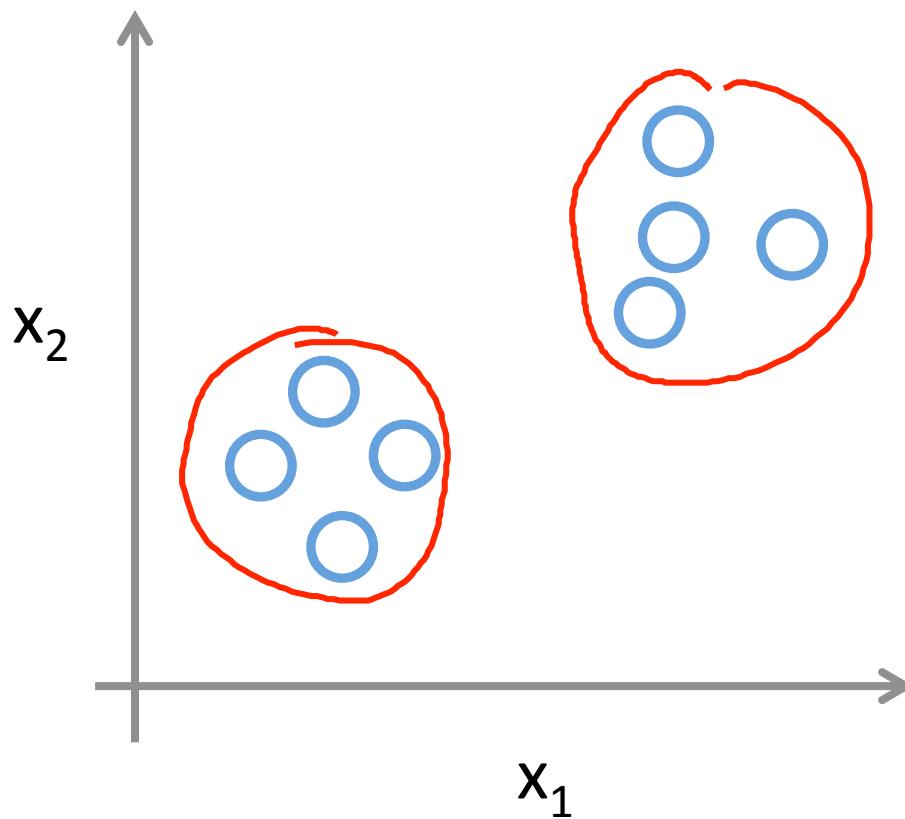
Introduction

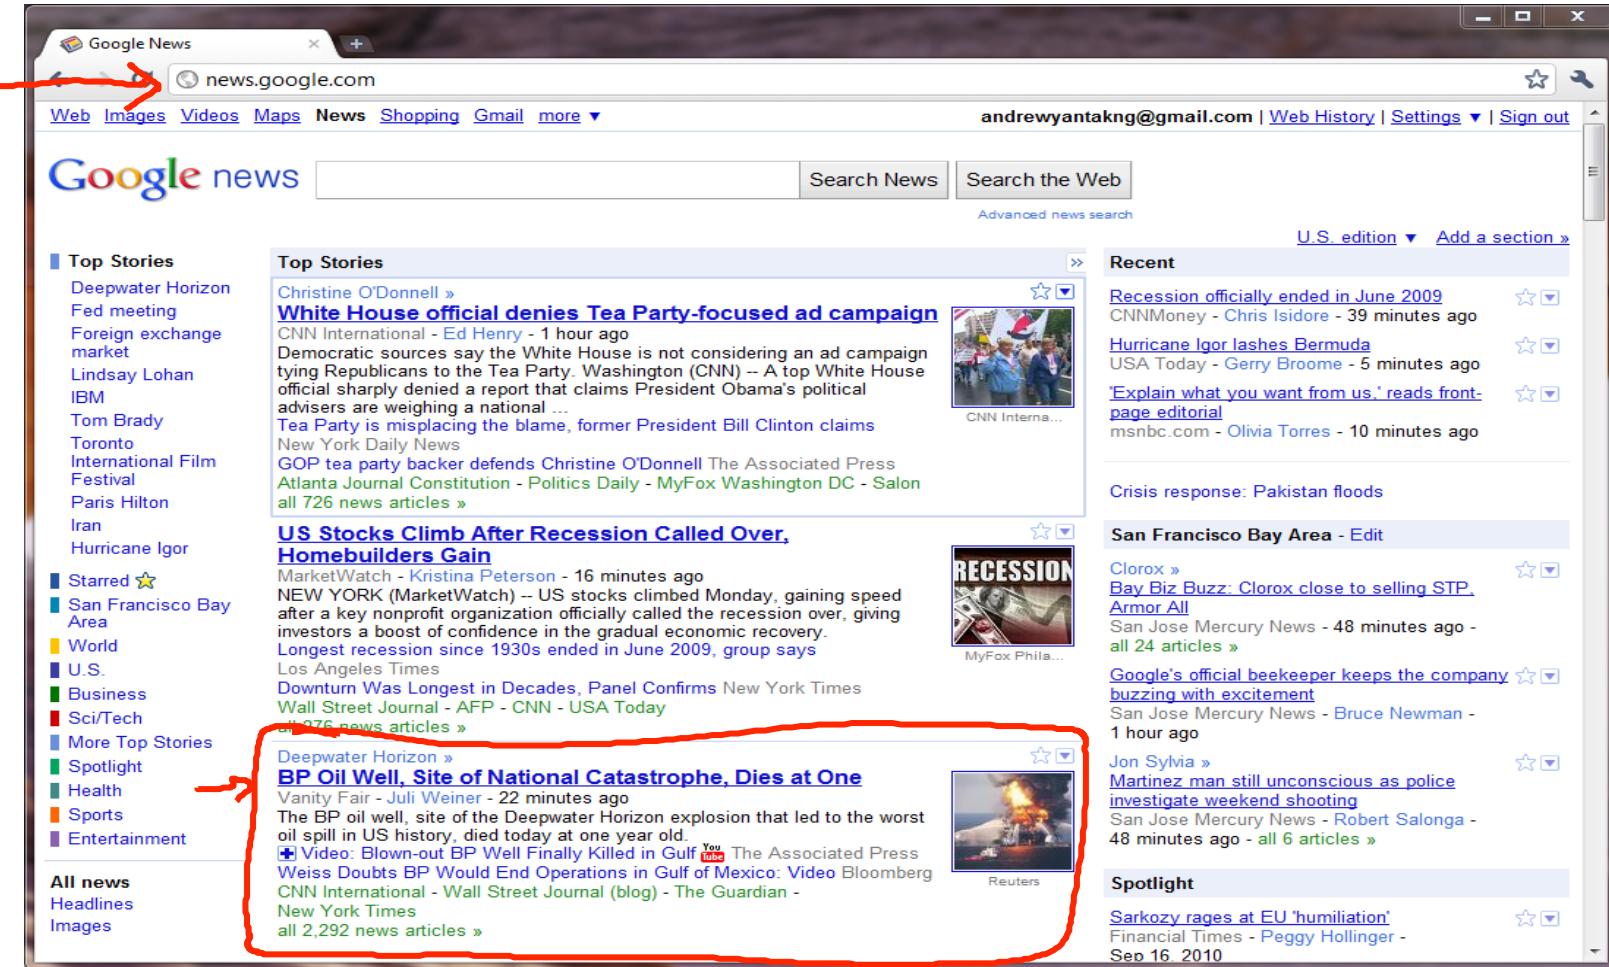
Unsupervised Learning

Supervised Learning



Unsupervised Learning





WSJ BP Kills Macondo, But Its ...

THE WALL STREET JOURNAL.
Digital Network WSJ.com MarketWatch BARRON'S All Things Digital. FINS SmartMoney More ▾

SEARCH

THE WALL STREET JOURNAL.

Log In Register For Free Subscribe Now, Get 2 Weeks Free

THE SOURCE

PREVIOUS NEXT

Financial Services Transport Leisure Insurance Oil & Gas Sport Caught on the Web Betting Technology

SEPTEMBER 20, 2010, 12:44 PM GMT

BP Kills Macondo, But Its Legacy Lives On

Article Comments (2)

Email Print Permalink Like 2 More Text

By James Herron

BP confirmed late Sunday that the Macondo well that leaked almost five million barrels of oil into the Gulf of Mexico has been permanently sealed, but the well will continue to affect BP and the wider oil industry for many years.

The most immediate worry for BP and its shareholders is how the authorities will apportion blame for the spill. BP's own investigation spread responsibility across



Associated Press

Fire boat response crews battled the blazing remnants of the off shore oil rig Deepwater Horizon on April 21, 2010.

Search The Source SEARCH

THE SOURCE HOME PAGE ▾

About The Source Follow Us:   

The Source is WSJ.com Europe's home for rapid-fire analysis of the day's big business and finance stories. It is edited by Lauren Mills, based in London.



Most Recent

Articles Comments

- Who Needs Plaza II Anyway
- Will Banks Be Forced to Split Retail And Banking Arms?
- Timing of Ratings Award Intriguing
- BP Kills Macondo, But Its Legacy Lives On
- We Already Need a Sequel to Basel III

CNN Allen: Well is dead, but m... edition.cnn.com/2010/US/09/20/gulf.oil.disaster/

Edition: INTERNATIONAL U.S. MÉXICO ARABIC
Set edition preference

CNN 

Sign up Log in SEARCH POWERED BY Google

Home Video World U.S. Africa Asia Europe Latin America Middle East Business World Sport Entertainment Tech Travel iReport

Allen: Well is dead, but much Gulf Coast work remains

By the CNN Wire Staff
September 20, 2010 -- Updated 1317 GMT (2117 HKT)



Click to play 

What next for Gulf oil spill?

STORY HIGHLIGHTS

(CNN) -- The ruptured Macondo well, a mile under the Gulf of Mexico off the Louisiana coast, has been pronounced dead.

Mixx Twitter Share Email Save Print

Recommend Be the first of your friends to recommend this.

Most Popular »

Today's five most popular stories

- White House official denies Tea Party-focused ad campaign
- Swedish far-right leader: Success due to immigration backlash
- Lady Gaga to lead rally against military's 'don't ask, don't tell'
- Why Facebook is blue -- six facts about Mark Zuckerberg
- China, Japan relations sour in dispute over fishing captain

[More](#)

BP oil spill cost hits near... www.guardian.co.uk/environment/2010/sep/20/bp-oil-spill-deepwater-horizon-costs-10bn

Mobile site Sign in Register A Text larger smaller About us Today's paper Zeitgeist

guardian.co.uk

Search Business Search

News | Sport | Comment | Culture | Business | Money | Life & style | Travel | Environment | TV | Video | Community | Blogs | Jobs

Business > BP

BP oil spill cost hits nearly \$10bn

BP has set up a \$20bn compensation fund after the Deepwater Horizon disaster, which has so far paid out 19,000 claims totalling more than \$240m

Julia Köllewe guardian.co.uk, Monday 20 September 2010 08.33 BST Article history



BP's costs for the Deepwater Horizon disaster have hit \$10bn. Photograph: Ho/Reuters

[bp](#) [Guardian](#) [Twitter](#) [Facebook](#) (3) [Tweet this \(68\)](#) [Comments \(10\)](#)

A larger | smaller

[Business](#) [BP](#)

[Environment](#) [BP oil spill](#) · [Oil](#) · [Oil spills](#)

[World news](#) [United States](#)

More news

More on this story



US government declares well 'dead'

BP oil spill news on Twitter

Latest news on the BP oil spill in the Gulf of Mexico

[bp](#) **BP_America:** Find fact sheets, open house info and e-mail alert sign-ups at @EPAgov's oil spill community page: <http://bit.ly/diuHoH> about 3 weeks ago

[guardianeco](#) UN report on Nigeria oil spills relies too much on data from Shell | Nnimmo Bassey <http://bit.ly/dBd7Ru> about 3 weeks, 5 days ago

[bp](#) **BP_America:** Newly discovered microbe thriving from consumption of #oil in the #Gulf of Mexico: <http://bit.ly/9kQYwa> about 3 weeks, 5 days ago

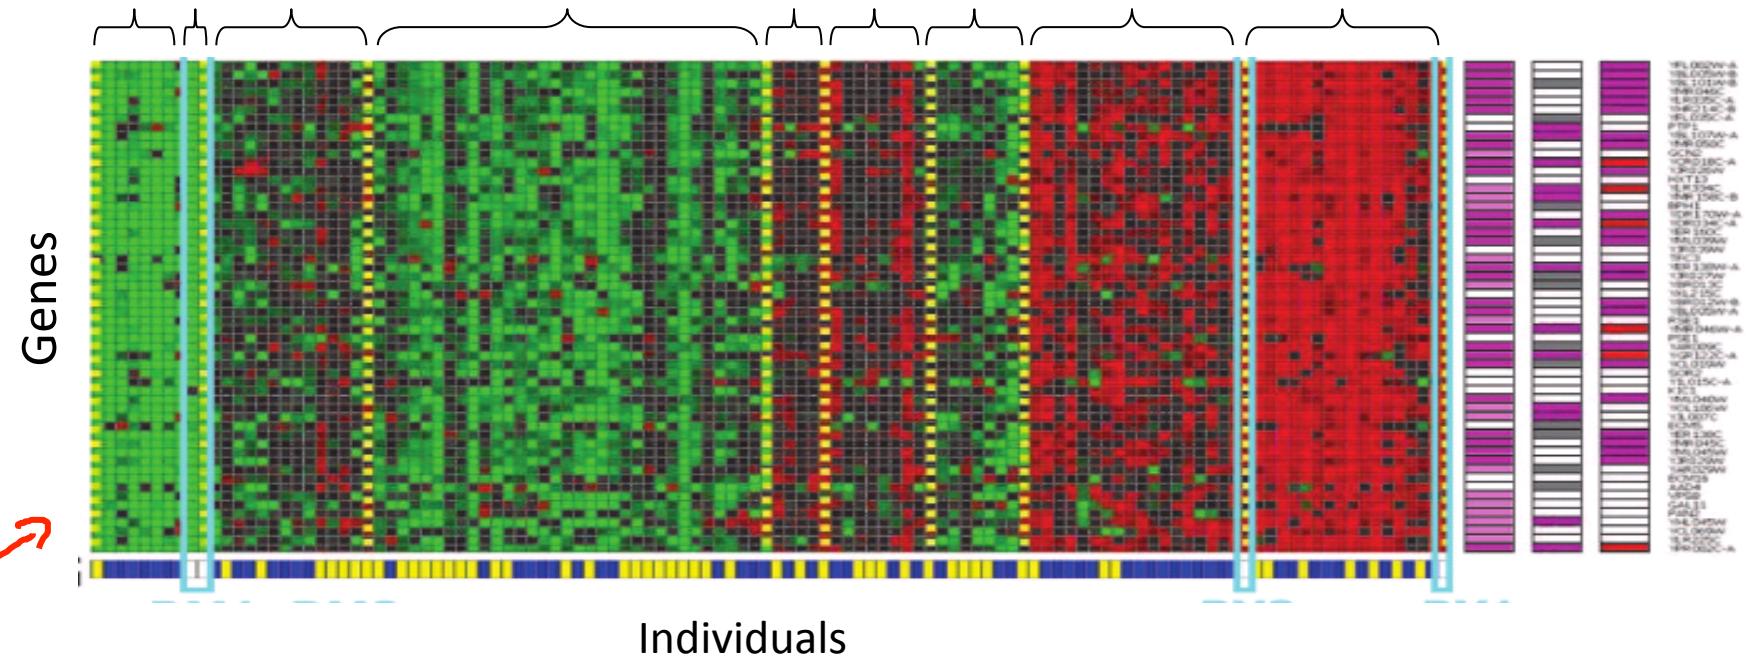
- Archive page: [read more BP oil spill tweets](#)
- Follow BP oil spill news with our Twitter list
- Full BP oil spill coverage from [guardian.co.uk](#)

On Business

Most viewed Zeitgeist Latest

Last 24 hours

1. Britain keeps AAA credit rating

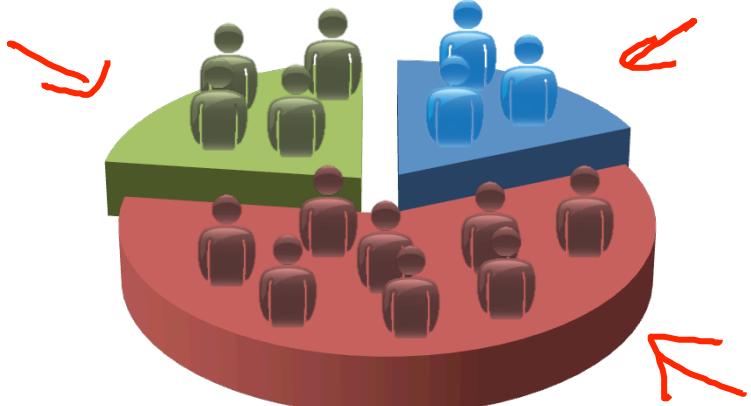


[Source: Su-In Lee, Dana Pe'er, Aimee Dudley, George Church, Daphne Koller]

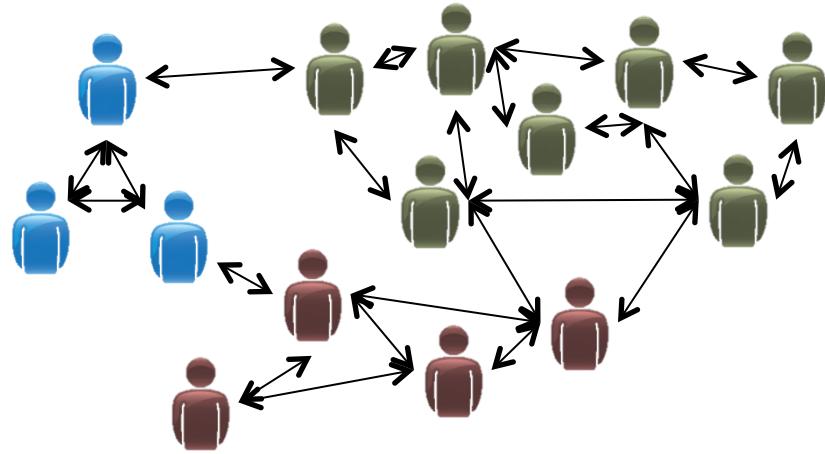
Andrew Ng



Organize computing clusters



Market segmentation



Social network analysis

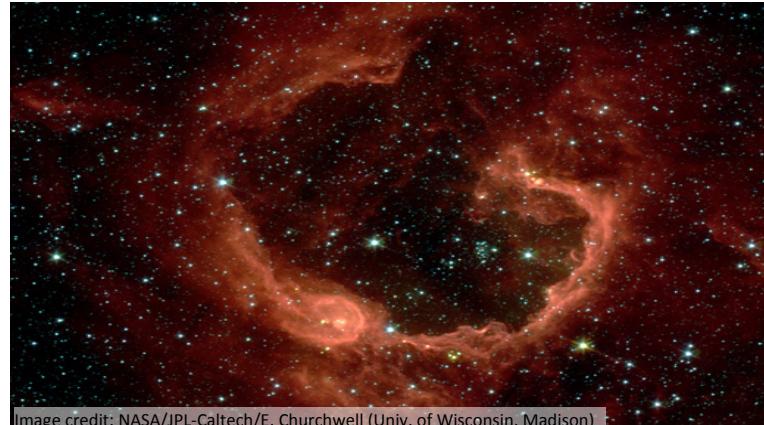
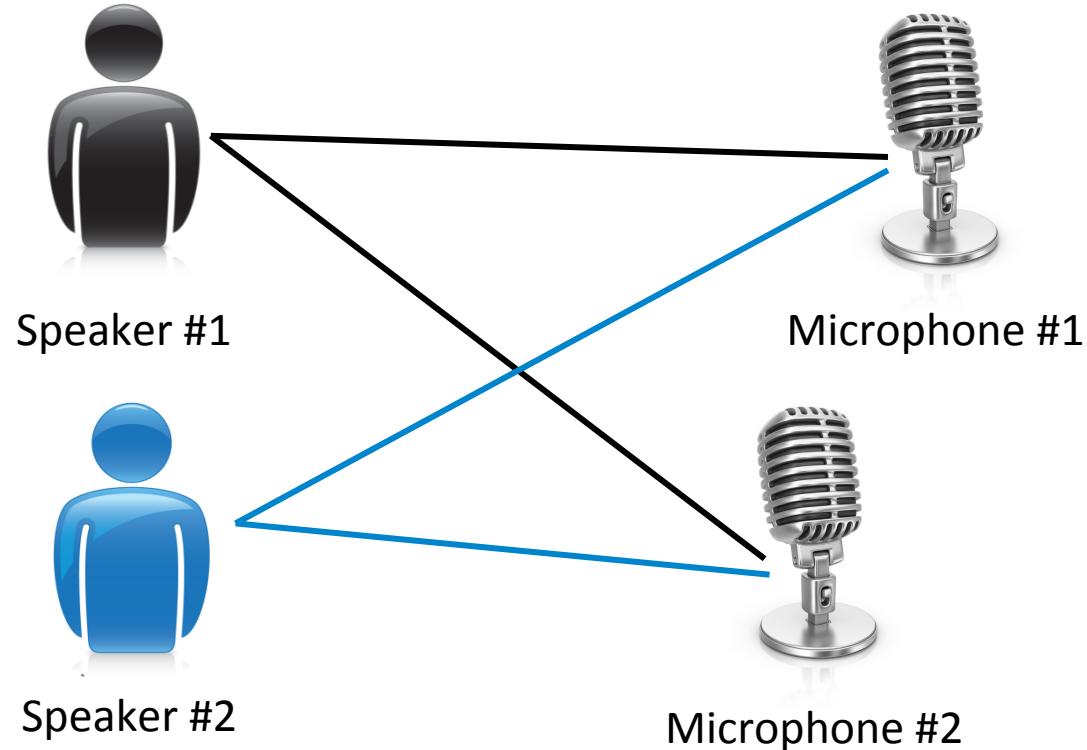


Image credit: NASA/JPL-Caltech/E. Churchwell (Univ. of Wisconsin, Madison)

Astronomical data analysis

Andrew Ng

Cocktail party problem



Microphone #1: 

Output #1: 

Microphone #2: 

Output #2: 

Microphone #1: 

Output #1: 

Microphone #2: 

Output #2: 

Cocktail party problem algorithm

```
[W,s,v] = svd((repmat(sum(x.*x,1),size(x,1),1).*x)*x');
```

Of the following examples, which would you address using an unsupervised learning algorithm? (Check all that apply.)

- Given email labeled as spam/not spam, learn a spam filter.
- Given a set of news articles found on the web, group them into
set of articles about the same story.
- Given a database of customer data, automatically discover market
segments and group customers into different market segments.
- Given a dataset of patients diagnosed as either having diabetes or
not, learn to classify new patients as having diabetes or not.